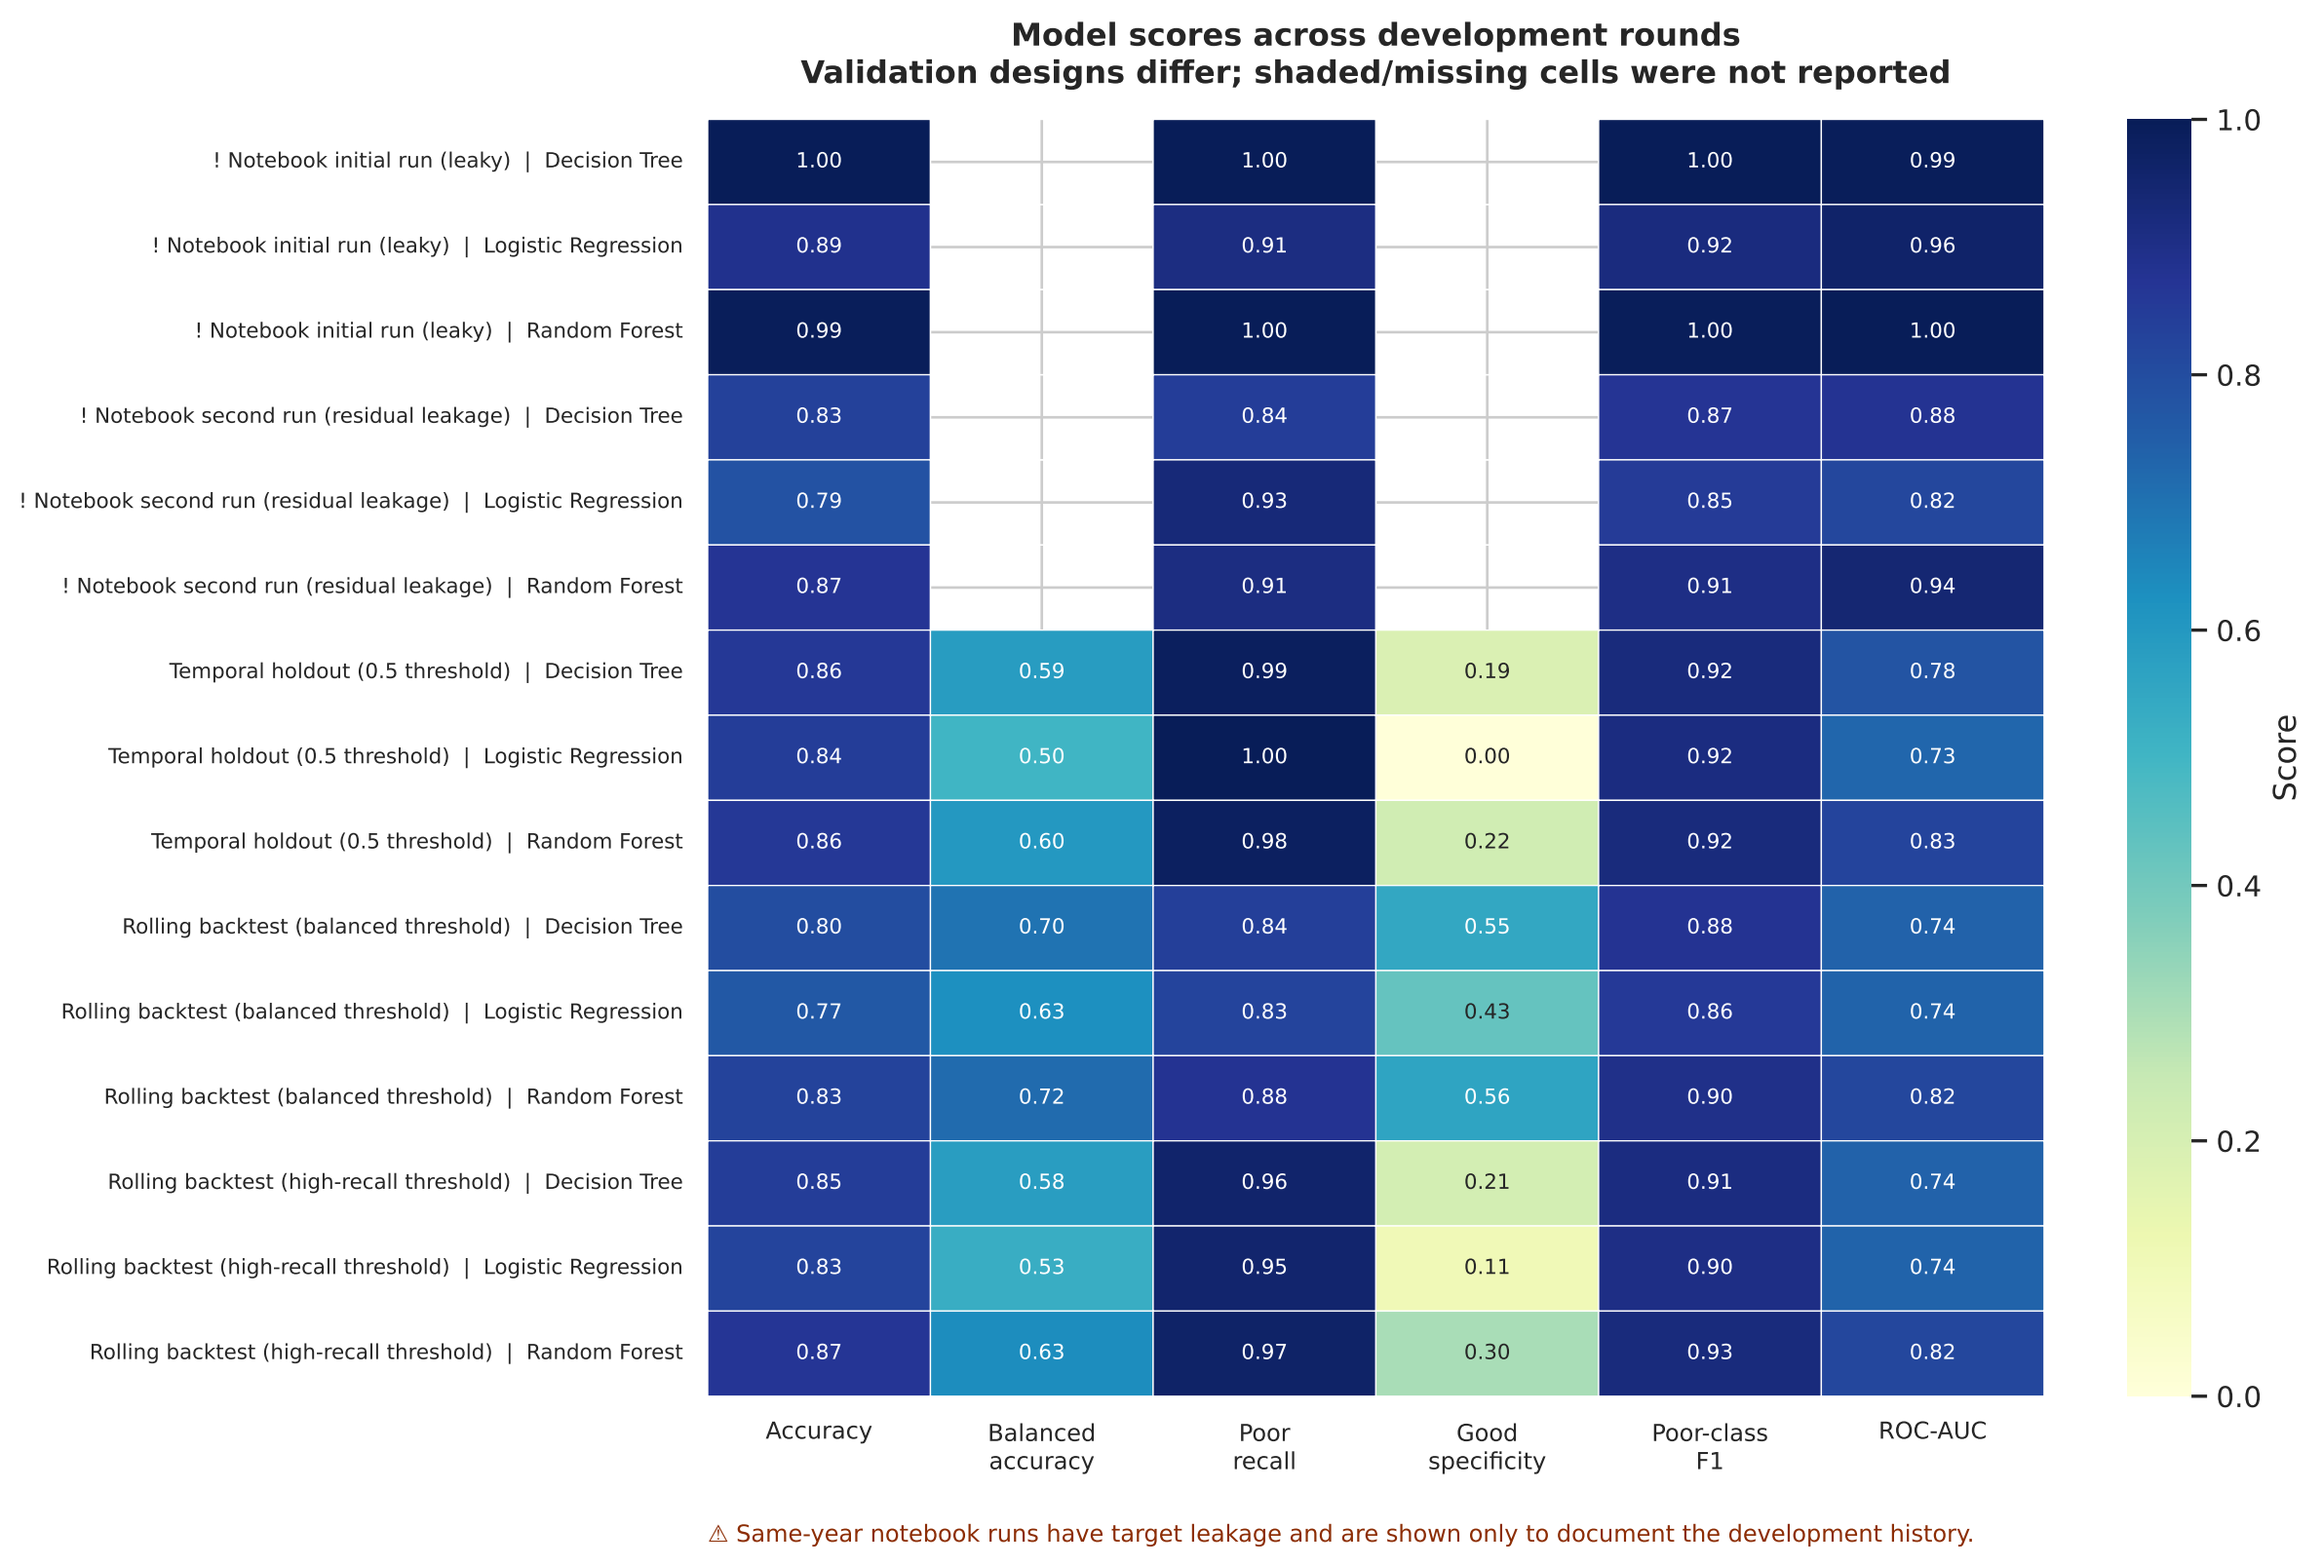
Model scores across development rounds
Validation designs differ; shaded/missing cells were not reported
! Notebook initial run (leaky) | Decision Tree
! Notebook initial run (leaky) | Logistic Regression
! Notebook initial run (leaky) | Random Forest
! Notebook second run (residual leakage) | Decision Tree
! Notebook second run (residual leakage) | Logistic Regression
! Notebook second run (residual leakage) | Random Forest
Temporal holdout (0.5 threshold) | Decision Tree
Temporal holdout (0.5 threshold) | Logistic Regression
Temporal holdout (0.5 threshold) | Random Forest
Rolling backtest (balanced threshold) | Decision Tree
Rolling backtest (balanced threshold) | Logistic Regression
Rolling backtest (balanced threshold) | Random Forest
Rolling backtest (high-recall threshold) | Decision Tree
Rolling backtest (high-recall threshold) | Logistic Regression
Rolling backtest (high-recall threshold) | Random Forest
1.00 1.00 1.00 0.99
0.89 0.91 0.92 0.96
0.99 1.00 1.00 1.00
0.83 0.84 0.87 0.88
0.79 0.93 0.85 0.82
0.87 0.91 0.91 0.94
0.86 0.59 0.99 0.19 0.92 0.78
0.84 0.50 1.00 0.00 0.92 0.73
0.86 0.60 0.98 0.22 0.92 0.83
0.80 0.70 0.84 0.55 0.88 0.74
0.77 0.63 0.83 0.43 0.86 0.74
0.83 0.72 0.88 0.56 0.90 0.82
0.85 0.58 0.96 0.21 0.91 0.74
0.83 0.53 0.95 0.11 0.90 0.74
0.87 0.63 0.97 0.30 0.93 0.82
Accuracy
Balanced accuracy
Poor recall
Good specificity
Poor-class F1
ROC-AUC
Score
1.0
0.8
0.6
0.4
0.2
0.0
⚠ Same-year notebook runs have target leakage and are shown only to document the development history.